

BEGUM ROKEYA UNIVERSITY, RANGPUR

DEPARTMENT OF STATISTICS

STATISTICAL DATA ANALYSIS

Comprehensive Laboratory Assignment Report

COURSE CODE: STAT4104

COURSE TITLE: CATEGORICAL DATA ANALYSIS

SUBMITTED BY

Name: MD OBAIDUL HOQUE

ID: 12110069

Reg No: 000015932

Session: 2021-2022

Year/Sem: 4th Year, 1st Semester

SUBMITTED TO

Name: DR. MD. SIDDIKUR RAHMAN

Associate Professor

Department of Statistics

Begum Rokeya University,
rangpur

SUBMISSION DATE: 18 MAY 2026

Problem 1

Consider the lung cancer data and solve the following questions using Python:

Basic EDA Questions

1. Summarize the distribution of Age. Is it normally distributed or skewed?
2. What proportion of individuals are smokers vs non-smokers?
3. Which income group has the highest frequency?
4. What percentage of individuals are exposed to high pollution?
5. What is the overall prevalence of lung disease?

Association & Crosstab

1. Is there an association between Smoking and Lung Disease?
2. Interpret using a contingency table.

Does Pollution Level Affect Lung Disease?

1. Compare lung disease prevalence across income groups.
2. Which factor appears to have the strongest relationship with lung disease?

Statistical Testing

1. Perform a Chi-square test for:
 - Smoking vs Lung Disease
 - Pollution vs Lung Disease
2. When should you use Fisher's Exact Test instead of Chi-square?
3. Calculate and interpret the Odds Ratio for Smoking and Lung Disease.

Correlation & Interpretation

1. Interpret the correlation between:

- Smoking and Lung Disease
 - Age and Lung Disease
2. Why is correlation not sufficient for causal inference?

Logistic Regression

1. Fit a logistic regression model:

[{LungDisease} ~ {Smoking} + {Age} + {Pollution} + {Income}]

2. Which variables are statistically significant predictors?
3. Interpret the odds ratio of Smoking.
4. What happens to the effect of Smoking after adjusting for Pollution?

Advanced Modeling

1. Add an interaction term:

[{Smoking} X {Pollution}]

2. Is the interaction significant?
3. Interpret the interaction.
4. Does pollution amplify smoking risk?

Model Evaluation

1. Compute the ROC curve and AUC.
2. Is the model good?
3. What is the confusion matrix?
4. Discuss model accuracy vs sensitivity vs specificity.

Research Question

“Using the lung disease dataset, perform exploratory data analysis and fit a logistic regression model to identify key determinants of lung disease. Interpret your findings and discuss public health implications.”

Basic EDA Questions

1. Summarize the distribution of Age. Is it normally distributed or skewed?

```
import pandas as pd
import numpy as np
from scipy.stats import chi2_contingency, fisher_exact
import statsmodels.formula.api as smf
from sklearn.metrics import roc_auc_score, confusion_matrix
from sklearn.metrics import accuracy_score, recall_score
df = pd.read_csv('lung_disease.csv')
print(df.head())
```

Outputs: print(df.head())

	ID	Smoking	Age	Income	Pollution	LungDisease
0	1	Yes	36	Low	High	Yes
1	2	No	28	Low	Low	No
2	3	No	52	Medium	High	No
3	4	No	39	High	High	No
4	5	Yes	32	Low	Low	Yes

In [5]:

```
# Summary statistics
print(df['Age'].describe())
```


```
count      500.000000
mean       43.904000 std
14.604841    min
20.000000    25%
30.750000
50%       45.000000 75%
56.000000 max     69.000000
Name: Age, dtype: float64
```

In [6]:

```
import seaborn as sns
import matplotlib.pyplot as plt
```


```
sns.histplot( df['Age'],
               kde=True, bins=10
             ) plt.title("Distribution of Age")

plt.xlabel("Age")
plt.ylabel("Frequency") plt.show()
```



In [7]:

```
sns.boxplot( x=df['Age']
             ) plt.title("Boxplot of Age")
plt.show()
```



In [8]:

```
#Skewness print(  
    "Skewness:", df['Age'].skew()  
)
```

Skewness: 0.012141455267137275

Normality Test

Using the Shapiro-Wilk Test

Null Hypothesis

[H_0 : {Age follows normal distribution}]

Alternative Hypothesis

[H_1 : {Age does not follow normal distribution}]

In [9]:

```
from scipy.stats import shapiro stat, p =  
shapiro(df['Age'])
```

```
print("Shapiro Test Statistic:", stat) print("P-value:", p)
```

Shapiro Test Statistic: 0.9488992946952997

P-value: 4.035055085974128e-12

In[10]:

```
if p > 0.05:  
    print("Age is approximately normally distributed")  
else:  
    print("Age is not normally distributed")
```

Age is not normally distributed

2. What proportion of individuals are smokers vs non-smokers?

In[11]:

```
# Count smokers and non-smokers  
smoking_counts =  
df['Smoking'].value_counts()  
print(smoking_counts)
```

Smoking

No 293

Yes 207

Name: count, dtype: int64

In[12]:

```
# Percentage of smokers and non-smokers  
  
smoking_percentage = ( df['Smoking'].value_counts(normalize=True)  
) * 100  
  
print(smoking_percentage)
```

Smoking

No 58.6

Yes 41.4

Name: proportion, dtype: float64

3. Which income group has the highest frequency?

```
In[14]: # Frequency of income groups income_counts =
df['Income'].value_counts()

print(income_counts)
# Most frequent income group

print(
    "Highest frequency income group is:",
    income_counts.idxmax() )
```

Income

Low 204

Medium 197

High 99

Name: count, dtype: int64

Highest frequency income group is: Low

4. What percentage of individuals are exposed to high pollution?

```
In[15]: # Percentage exposed to high pollution

high_pollution_percentage = (
    (df['Pollution'] == 'High').mean()
) * 100

print(
    "Percentage exposed to high pollution:", high_pollution_percentage
)
```

Percentage exposed to high pollution: 70.39999999999999 5. What is the

```
overall # Overall prevalence of lung disease

lung_disease_percentage = (
prevalence (df['LungDisease'] == 'Yes').mean()
) * 100

of lung print(
    "Overall prevalence of lung disease:", lung_disease_percentage
)
```


disease? In[16]:

Overall prevalence of lung disease: 52.800000000000004

Association & Crosstab

1. Is there an association between Smoking and Lung Disease?

In[17]:

```
# Contingency table  
  
table = pd.crosstab( df['Smoking'],  
                    df['LungDisease']  
                    ) print(table)
```

LungDisease	No	Yes
Smoking		
No	184	109

Yes 52 155 In[18]:

```
# Row percentage table

row_percent = pd.crosstab( df['Smoking'],
                           df['LungDisease'],
                           normalize='index'
                           ) * 100

print(row_percent.round(2))
```


LungDisease	No	Yes
Smoking		
No	62.80	37.20
Yes	25.12	74.88

In[19]:

```
import seaborn as sns import
matplotlib.pyplot as plt

sns.heatmap( table,
             annot=True,
             fmt='d',
             cmap='Blues'
             )

plt.title(
    "Smoking vs Lung Disease"
) plt.show()
```



2. Interpret using a contingency table.

Interpretation

The contingency table shows that smokers have a much higher proportion of lung disease compared to non-smokers.

Among smokers, approximately **74.88%** have lung disease, while only 37.20% of non-smokers have lung disease.

This suggests a strong positive association between smoking and lung disease.

Does Pollution Level Affect Lung Disease?

1. Compare lung disease prevalence across income groups.

In[23]:

```
# Income vs Lung Disease

income_table = pd.crosstab(
    df['Income'], df['LungDisease']
)

print(income_table)
```

LungDisease	No	Yes
Income		
High	51	48
Low	90	114
Medium	95	102

In[24]:

```
# Row percentage table

income_percent = pd.crosstab(
    df['Income'], df['LungDisease'],
    normalize='index'
) * 100

print(income_percent.round(2))
```

LungDisease	No	Yes
Income		
High	51.52	48.48
Low	44.12	55.88
Medium	48.22	51.78

Interpretation

The prevalence of lung disease varies across income groups.

Individuals in the **Low income** group have the highest prevalence of lung disease, with approximately **55.88%** affected.

The **Medium income** group has a lung disease prevalence of 51.78%, while the **High income** group has the lowest prevalence at **48.48%**.

This suggests that lower income groups may be more affected by lung disease compared to higher income groups.

2. Which factor appears to have the strongest relationship with lung disease?

In[25]:

```
print("Smoking vs Lung Disease") print(row_percent.round(2))

print("\nPollution vs Lung Disease") print(pollution_percent.round(2))

print("\nIncome vs Lung Disease") print(income_percent.round(2))
```

Smoking vs Lung Disease

LungDisease	No	Yes
-------------	----	-----

Smoking		
---------	--	--

No	62.80	37.20
----	-------	-------

Yes	25.12	74.88
-----	-------	-------

Pollution vs Lung Disease

LungDisease	No	Yes
-------------	----	-----

Pollution		
-----------	--	--

High	38.64	61.36
------	-------	-------

Low	67.57	32.43
-----	-------	-------

Income vs Lung Disease

LungDisease	No	Yes
-------------	----	-----

Income		
--------	--	--

High	51.52	48.48
------	-------	-------

Low	44.12	55.88
-----	-------	-------

Medium	48.22	51.78
--------	-------	-------

Interpretation

Among the examined factors, **Smoking** appears to have the strongest relationship with lung disease.

Approximately **74.88%** of smokers have lung disease compared to only 37.20% of non-smokers, showing a large difference between groups.

Pollution also shows a noticeable association, as individuals exposed to high pollution have a higher prevalence of lung disease (61.36%) compared to those exposed to low pollution (**32.43%**).

Income shows comparatively smaller differences in lung disease prevalence across groups.

Statistical Testing

1. Perform a Chi-square test for:

- Smoking vs Lung Disease

- Pollution vs Lung Disease

In[26]:

```
#Smoking vs Lung Disease — Chi-Square Test from
scipy.stats import chi2_contingency

# Contingency table

smoking_table = pd.crosstab( df['Smoking'],
                             df['LungDisease']
                           ) print(smoking_table)
```

LungDisease	No	Yes
Smoking		
No	184	109
Yes	52	155

In[27]:

```
# Chi-square test

chi2, p, dof, expected = chi2_contingency( smoking_table
) print("Chi-square Statistic:", chi2)
print("Degrees of Freedom:", dof) print("P-
value:", p)
```

Chi-square Statistic: 67.5942649754021
Degrees of Freedom: 1
P-value: 2.0085094601569646e-16

Interpretation

The Chi-square test results indicate a statistically significant association between Smoking and Lung Disease.

Since the p-value is much smaller than 0.05, we reject the null hypothesis of independence.

This suggests that smoking status is significantly related to lung disease occurrence.

```
In[29]: #Pollution vs Lung Disease — Chi-Square Test
# Contingency table

pollution_table = pd.crosstab( df['Pollution'],
                                df['LungDisease']
                                ) print(pollution_table)
```

```
LungDisease  No  Yes
Pollution
High         136 216
Low          100  48
```

```
In [30]: # Chi-square test

chi2, p, dof, expected = chi2_contingency( pollution_table
) print("Chi-square Statistic:", chi2)
print("Degrees of Freedom:", dof) print("P-
value:", p)
```

```
Chi-square Statistic: 33.84256825473482
Degrees of Freedom: 1
P-value: 5.975692575712951e-09
```

Interpretation

The Chi-square test results show a statistically significant association between Pollution Level and Lung Disease.

Since the p-value is less than 0.05, we reject the null hypothesis of independence.

This suggests that pollution level is significantly related to the occurrence of lung disease.

2. When Should Fisher's Exact Test Be Used?

Fisher's Exact Test should be used instead of the Chi-square test when sample size is small or when expected cell frequencies are less than 5 in a contingency table.

3. Odds Ratio for Smoking and Lung Disease

```
In[31]: print(smoking_table)
```

LungDisease	No	Yes
Smoking		
No	184	109
Yes	52	155

```
In[32]: #Calculate Odds Ratio
from statsmodels.stats.contingency_tables import Table2x2

# Convert to numpy array table_array =

smoking_table.values

# Odds ratio oddsratio =
Table2x2(
    table_array
)

print(
    "Odds Ratio:", oddsratio.oddsratio
)
```

Odds Ratio: 5.031757233592096

Interpretation

The odds ratio is approximately 5.03, which indicates that smokers have about 5 times higher odds of developing lung disease compared to non-smokers.

This suggests a strong positive association between smoking and lung disease.

Correlation & Interpretation

1. Interpret the correlation between:
 - Smoking and Lung Disease
 - Age and Lung Disease


```
In[33]: # Convert categorical variables to numeric

df['Smoking_num'] = df['Smoking'].map({
    'No': 0,
    'Yes': 1
})

df['LungDisease_num'] = df['LungDisease'].map({
    'No': 0,
    'Yes': 1 })
```

```
In[34]: # Correlation between Smoking and Lung Disease

correlation_smoking = df[
    ['Smoking_num', 'LungDisease_num']
].corr() print(correlation_smoking)
```

	Smoking_num	LungDisease_num
Smoking_num	1.000000	0.371747
LungDisease_num	0.371747	1.000000

Interpretation

The correlation coefficient between Smoking and Lung Disease is 0.372, indicating a moderate positive relationship.

This suggests that individuals who smoke are more likely to develop lung disease compared to non-smokers.

2. Why is correlation not sufficient for causal inference?

Correlation is not sufficient for causal inference because correlation only measures the strength and direction of association between variables.

A correlation between two variables does not necessarily mean that one variable causes the other.

Other factors, called confounding variables, may influence both variables simultaneously. In addition, correlation cannot establish temporal order or rule out coincidence.

Therefore, further statistical analysis and experimental or longitudinal studies are needed to establish causation.

Logistic Regression

1. Fit a logistic regression model:

[{LungDisease} ~ {Smoking} + {Age} + {Pollution} + {Income}]

2. Which variables are statistically significant predictors?
3. Interpret the odds ratio of Smoking.
4. What happens to the effect of Smoking after adjusting for Pollution?

1. Fit a logistic regression model:

[{LungDisease} ~ {Smoking} + {Age} + {Pollution} + {Income}]

In[35]: `import statsmodels.formula.api as smf`

In[36]: `# Fit logistic regression model model =
smf.logit(

 'LungDisease_num ~ Smoking_num + Age + C(Pollution) + C(Income)', data=df
).fit()`

Optimization terminated successfully. Current
function value: 0.552949
Iterations 6

In[37]: `# View model summary print(model.summary())`

Logit Regression Results

```
=====
Dep. Variable:    LungDisease_num  No. Observations:      500
Model:            Logit  Df Residuals:      494
Method:           MLE    Df Model:          5
Date:            Sat, 16 May 2026  Pseudo R-squ.:    0.2005 Time:      09:35:59
Log-Likelihood:    -276.47 converged:      True LL-Null:      -345.79
Covariance Type:  nonrobust    LLR p-value:      3.499e-28
=====
```

coef	std err	z	P> z	[0.025	0.975]

Intercept	-2.1940	0.425	-5.164	0.000	-3.027

-1.361						
C(Pollution)[T.Low]	-1.3027	0.235	-5.533	0.000	-1.764	
-0.841						
C(Income)[T.Low]	0.4396	0.285	1.544	0.123	-0.119	
0.998 C(Income)[T.Medium]	0.2201	0.285	0.773	0.440	-0.338	
0.779 Smoking_num	1.7022	0.218	7.801	0.000	1.275	
2.130 Age	0.0399	0.007	5.455	0.000	0.026	
0.054						

=====

=====

=====

Interpretation

The logistic regression model is statistically significant, indicating that the predictors collectively explain lung disease occurrence.

Significant Predictors

- Smoking is a statistically significant predictor of lung disease (($p < 0.001$)).
- Age is also statistically significant (($p < 0.001$)).
- Pollution level significantly affects lung disease (($p < 0.001$)).

Income

The income variables are not statistically significant because their p-values are greater than 0.05.

Smoking Effect

The positive coefficient for Smoking indicates that smokers have a higher probability of developing lung disease compared to non-smokers.

Age Effect

The positive coefficient for Age suggests that the probability of lung disease increases with age.

Pollution Effect

Individuals exposed to low pollution have lower odds of lung disease compared to those exposed to high pollution.

2. Which variables are statistically significant predictors?

In [38]:

```
# Extract p-values print(model.pvalues)
```

```
Intercept          2.413178e-07
C(Pollution)[T.Low] 3.139044e-08
C(Income)[T.Low]    1.226431e-01
C(Income)[T.Medium] 4.397235e-01
Smoking_num        6.131827e-15 Age
4.906629e-08 dtype: float64
```

Interpretation

The statistically significant predictors of lung disease are:

- **Smoking** (($p < 0.001$))
- **Age** (($p < 0.001$))
- **Pollution** (($p < 0.001$))

These variables have p-values less than 0.05, indicating a significant association with lung disease.

The income variables are not statistically significant because their p-values are greater than 0.05.

3. Interpret the Odds Ratio of Smoking

In [39]:

```
import numpy as np #
Odds ratios
odds_ratios = np.exp( model.params
) print(odds_ratios)
```

```
Intercept          0.111471
C(Pollution)[T.Low] 0.271784
C(Income)[T.Low]    1.552101
C(Income)[T.Medium] 1.246243
Smoking_num        5.485739 Age
1.040683 dtype: float64
```

```
In [40]: print(
    "Odds Ratio of Smoking:",
    odds_ratios['Smoking_num'] )
```

Odds Ratio of Smoking: 5.485738801540751

Interpretation

The odds ratio for Smoking is approximately 5.49.

This indicates that smokers have about 5.5 times higher odds of developing lung disease compared to non-smokers, after adjusting for age, pollution, and income.

4. What happens to the effect of Smoking after adjusting for Pollution?

```
In [41]: # Logistic regression without Pollution

model_without_pollution = smf.logit(
    'LungDisease_num ~ Smoking_num + Age + C(Income)',
    data=df
).fit()

print(model_without_pollution.summary())
```

Optimization terminated successfully. Current
function value: 0.585871
Iterations 5

Logit Regression Results

```
=====
Dep. Variable:    LungDisease_num  No. Observations:      500
Model:            Logit  Df Residuals:      495
Method:           MLE  Df Model:           4
Date:            Sat, 16 May 2026  Pseudo R-squ.:    0.1528 Time:      09:41:00
Log-Likelihood:   -292.94 converged:      True  LL-Null:      -345.79
Covariance Type:      nonrobust    LLR p-value:      5.986e-22
=====
```

```
coef    std err      z    P>|z|    [0.025
0.975]
```

```
-----
Intercept          -2.5652    0.418   -6.138    0.000   -3.384
-1.746
C(Income)[T.Low]    0.4658    0.278    1.678    0.093   -0.078
1.010 C(Income)[T.Medium] 0.2646    0.278    0.952    0.341   -0.280
```

0.809 Smoking_num	1.6651	0.210	7.938	0.000	1.254
2.076 Age	0.0395	0.007	5.551	0.000	0.026
0.053					

In [42]:

```
print(
    "Smoking coefficient without Pollution:",
    model_without_pollution.params['Smoking_num']
)

print(
    "Smoking coefficient with Pollution:",
    model.params['Smoking_num'] )
```

Smoking coefficient without Pollution: 1.6651027434859365

Smoking coefficient with Pollution: 1.7021517795592478

In [43]:

```
# Odds ratios

or_without_pollution = np.exp( model_without_pollution.params['Smoking_num']
)

or_with_pollution = np.exp( model.params['Smoking_num']

)

print(
    "Odds Ratio without Pollution:", or_without_pollution
)

print(
    "Odds Ratio with Pollution:", or_with_pollution
)
```

Odds Ratio without Pollution: 5.286216346155295

Odds Ratio with Pollution: 5.485738801540751

Interpretation

The effect of Smoking remains strong even after adjusting for Pollution.

The smoking coefficient and odds ratio changed only slightly after including pollution in the model:

- Odds Ratio without Pollution = 5.29

- Odds Ratio with Pollution = 5.49

This suggests that smoking independently contributes to lung disease risk, and its effect is not fully explained by pollution exposure.

Advanced Modeling

1. Add an interaction term:

[{Smoking} X {Pollution}]

2. Is the interaction significant?
3. Interpret the interaction.
4. Does pollution amplify smoking risk?

1. Add an interaction term:

[{Smoking} X {Pollution}]

In [45]:

```
# Logistic regression with interaction term interaction_model =
smf.logit(

    'LungDisease_num ~ Smoking_num * C(Pollution) + Age + C(Income)', data=df
).fit()
print(interaction_model.summary())
```

Optimization terminated successfully. Current
function value: 0.552178
Iterations 6

Logit Regression Results

```
=====
Dep. Variable:    LungDisease_num  No. Observations:      500
Model:            Logit  Df Residuals:      493
Method:           MLE  Df Model:           6
Date:            Sat, 16 May 2026  Pseudo R-squ.:      0.2016 Time:            09:44:44
Log-Likelihood:    -276.09 converged:      True  LL-Null:            -
345.79 Covariance Type:      nonrobust  LLR p-value:      1.341e-27
=====

coef  std err      z    P>|z|
[0.025  0.975]
```

Intercept	-2.1285	0.430	-4.953	0.000				
-2.971	-1.286							
C(Pollution)[T.Low]	-1.4883	0.323	-4.606	0.000				
-2.122	-0.855							
C(Income)[T.Low]	0.4465	0.285	1.567	0.117				
-0.112	1.005							
C(Income)[T.Medium]	0.2214	0.285	0.777	0.437				
-0.337	0.780							
Smoking_num	1.5778	0.257	6.130	0.000				
1.073	2.082	Smoking_num:C(Pollution)[T.Low]	0.4217	0.483	0.874	0.382		
-0.524	1.367							
Age	0.0393	0.007	5.352	0.000				
0.025								0.054

2. Is the interaction significant?

The interaction term between Smoking and Pollution is not statistically significant because:

[$p = 0.382 > 0.05$]

This indicates that there is insufficient evidence of a significant interaction effect between smoking and pollution on lung disease risk.

3. Interpret the interaction

The positive interaction coefficient suggests that the combined effect of smoking and pollution may increase lung disease risk more than their individual effects alone.

However, since the interaction term is not statistically significant, this additional combined effect is not strong enough to conclude a meaningful interaction in the population.

4. Does pollution amplify smoking risk?

The interaction coefficient is positive, which suggests that pollution may increase the effect of smoking on lung disease risk.

However, because the interaction term is not statistically significant ($(p > 0.05)$), there is no strong statistical evidence that pollution significantly amplifies smoking risk.

Model Evaluation

1. Compute the ROC curve and AUC.
2. Is the model good?

3. What is the confusion matrix?
 4. Discuss model accuracy vs sensitivity vs specificity.
-
1. Compute the ROC curve and AUC.

In [46]: *# Predicted probabilities* y_pred_prob =
model.predict(df)

In [47]: *#ROC curve and AUC* **from**
sklearn.metrics **import** (
roc_curve, roc_auc_score
)

ROC components

fpr, tpr, thresholds = roc_curve(
df['LungDisease_num'], y_pred_prob
) *# AUC*

score

auc = roc_auc_score(
df['LungDisease_num'], y_pred_prob
) print("AUC:", auc)

AUC: 0.7871645480225988


In [48]:

```
import matplotlib.pyplot as plt

plt.plot(fpr,
         tpr,
         label=f'AUC = {auc:.3f}')

plt.plot(
    [0, 1], [0, 1],
    linestyle='--')

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve") plt.legend()
plt.show()
```



2. Is the model good?

Interpretation

The model appears to be reasonably good because the AUC value is approximately 0.787.

An AUC between 0.7 and 0.8 indicates acceptable discrimination ability, meaning the model can fairly well distinguish between individuals with and without lung disease.

Therefore, the logistic regression model has a good predictive performance, although there is still room for improvement.

3. What is the confusion matrix?

In [49]:

```
# Predicted classes  
y_pred_class = (y_pred_prob  
                >= 0.5
```

```
).astype(int)
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(
    df['LungDisease_num'], y_pred_class
) print(cm)
```


```
[[156  80]
 [ 64 200]] In
```

[50]:

```
import seaborn as sns

sns.heatmap( cm,
    annot=True,
    fmt='d',
    cmap='Blues'
)

plt.xlabel("Predicted") plt.ylabel("Actual")
plt.title("Confusion Matrix") plt.show()
```



Interpretation of Confusion Matrix

The confusion matrix shows that:

- 156 individuals without lung disease were correctly classified (True Negatives).
- 200 individuals with lung disease were correctly classified (True Positives).
- 80 individuals without lung disease were incorrectly classified as having lung disease (False Positives).
- 64 individuals with lung disease were incorrectly classified as not having lung disease (False Negatives).

Overall, the model correctly classifies a large number of cases, indicating reasonably good predictive performance.

4. Discuss model accuracy vs sensitivity vs specificity.

In [51]: `from sklearn.metrics import ( accuracy_score,`

```

    recall_score
) #

Accuracy

accuracy = accuracy_score(
    df['LungDisease_num'], y_pred_class
) # Sensitivity (Recall)

sensitivity = recall_score(
    df['LungDisease_num'], y_pred_class
)

# Specificity tn, fp, fn, tp =

cm.ravel() specificity = tn / (tn +

fp)

print("Accuracy:", accuracy) print("Sensitivity:",
sensitivity) print("Specificity:", specificity)

```

Accuracy: 0.712

Sensitivity: 0.7575757575757576

Specificity: 0.6610169491525424

Interpretation

The model achieved an accuracy of approximately 71.2%, meaning that about 71% of the observations were correctly classified.

The sensitivity is approximately 75.8%, indicating that the model is fairly good at correctly identifying individuals with lung disease.

The specificity is approximately 66.1%, meaning the model is moderately effective at correctly identifying individuals without lung disease.

Overall, the model performs reasonably well, with better sensitivity than specificity, suggesting it is more effective at detecting lung disease cases than non-disease cases.

Research Question

“Using the lung disease dataset, perform exploratory data analysis and fit a logistic regression model to identify key determinants of lung disease. Interpret your findings and discuss public health implications.”

Research Question

Exploratory Data Analysis and Logistic Regression Findings

Exploratory data analysis showed noticeable differences in lung disease prevalence across smoking status, pollution exposure, and income groups.

Smokers had a substantially higher prevalence of lung disease compared to nonsmokers. Individuals exposed to high pollution also showed a higher proportion of lung disease cases. Lower income groups appeared to have slightly higher lung disease prevalence than higher income groups.

Chi-square tests indicated statistically significant associations between:

- Smoking and Lung Disease
- Pollution and Lung Disease

The odds ratio analysis showed that smokers had approximately 5 times higher odds of developing lung disease compared to non-smokers.

A logistic regression model was fitted using Smoking, Age, Pollution, and Income as predictors of lung disease.

The model identified:

- Smoking
- Age
- Pollution

as statistically significant predictors of lung disease, while income variables were not statistically significant after adjustment.

The logistic regression results indicated that:

- smokers had significantly higher odds of lung disease,
- increasing age was associated with higher lung disease risk,
- and exposure to high pollution increased the likelihood of lung disease.

The interaction analysis between Smoking and Pollution suggested a positive interaction effect, although the interaction term was not statistically significant.

Model evaluation showed:

- $AUC \approx 0.787$
- Accuracy $\approx 71.2\%$
- Sensitivity $\approx 75.8\%$
- Specificity $\approx 66.1\%$

These results indicate that the model has acceptable predictive performance and is reasonably effective at identifying individuals with lung disease.

Public Health Implications

The findings highlight smoking and pollution exposure as major risk factors for lung disease.

Public health interventions should focus on:

- reducing smoking prevalence,
- improving air quality,
- increasing awareness regarding respiratory health,
- and promoting early screening for high-risk populations.

Policies targeting smoking cessation and environmental pollution control may substantially reduce the burden of lung disease in the population.

In []:

Problem-2

Suppose we want to determine if three different exercise programs impact weight loss differently. We recruit 90 people to participate in an experiment in which we randomly assign (using uniform distribution) 30 people to follow either Program A, Program B, or Program C for one month. The predictor variable being studied is the exercise program, and the response variable is weight loss, measured in pounds.

Conduct a one-way ANOVA to determine if there is a statistically significant difference between the resulting weight loss from the three programs, check the model assumptions, and analyze treatment differences.

In [30]:

```
#Import libraries import
pandas as pd import
numpy as np

import matplotlib.pyplot as plt

import seaborn as sns from scipy

import stats

import statsmodels.api as sm from statsmodels.formula.api import
ols from statsmodels.stats.multicomp import pairwise_tukeyhsd
```

In [31]:

```
#Create dataset np.random.seed(42)

A = np.random.normal(loc=5, scale=1.5, size=30)
B = np.random.normal(loc=8, scale=1.5, size=30)
C = np.random.normal(loc=4, scale=1.5, size=30)
```

In [32]:

```
#Combine into DataFrame df
= pd.DataFrame({
    'WeightLoss': np.concatenate([A, B, C]),
    'Program': ['A']*30 + ['B']*30 + ['C']*30
}) df.head(35)
```

Out[32]:

WeightLoss Program

0	5.745071	A
1	4.792604	A
2	5.971533	A
3	7.284545	A
4	4.648770	A
5	4.648795	A
6	7.368819	A
7	6.151152	A
8	4.295788	A
9	5.813840	A
10	4.304873	A

11	4.301405	A
12	5.362943	A
13	2.130080	A
14	2.412623	A
15	4.156569	A
16	3.480753	A
17	5.471371	A
18	3.637964	A
19	2.881544	A
20	7.198473	A
21	4.661336	A
22	5.101292	A
23	2.862878	A
24	4.183426	A
25	5.166384	A
26	3.273510	A
27	5.563547	A
28	4.099042	A
29	4.562459	A
30	7.097440	B
WeightLoss Program		
31	10.778417	B
32	7.979754	B
33	6.413434	B
34	9.233817	B

In [33]:

```

###Basic EDA##
#Group Means
df.groupby('Program')['WeightLoss'].mean()

```

Out[33]:

WeightLoss	
Program	
A	4.717780
B	7.818256
C	4.019327

dtype: float64 In

[34]:

```
#Group Standard Deviations  
df.groupby('Program')['WeightLoss'].std()
```


Out[34]:

WeightLoss	
Program	
A	1.350010
B	1.396653
C	1.487975

dtype: float64

In [35]:

```
#Boxplot  
sns.boxplot(x='Program', y='WeightLoss', data=df)  
plt.title("Weight Loss by Exercise Program") plt.show()
```




The boxplot shows that weight loss differs across the three exercise programs.

Program B has the highest median weight loss, indicating it is the most effective program overall. Program A shows moderate weight loss with less variability. Program C has the lowest median weight loss and includes a possible outlier near 0.

In [36]:

```
#Distribution Plot  
sns.histplot(data=df, x='WeightLoss', hue='Program', kde=True) plt.show()
```



ONE way ANOVA test

ANOVA Assumptions

Before ANOVA we check:

1. Independence
2. Normality
3. Equal Variance

```
In [37]: #Normality check(Shapiro-wilk test) for
group in df['Program'].unique():

    stat, p = stats.shapiro( df[df['Program']==group]['WeightLoss']
    )

    print(group) print("P-value:", p)
    print()
```

A
P-value: 0.686805494291708

B

P-value: 0.9129582559088034

C

P-value: 0.36537361581254685

Since $p > 0.05$ so it is normal

In [38]:

```
#Homogeneity of Variance  
  
#Levene's Test stat, p =  
stats.levene(A, B, C) print("P-  
value:", p)
```

P-value: 0.8626638812717601

Since $p > 0.05$ so equal variances

In [40]:

```
del C
```

In [41]:

```
model = ols(  
    'WeightLoss ~ C(Program)',  
    data=df  
) .fit()
```

In [42]:

```
#ANOVA Table  
anova_table = sm.stats.anova_lm(model, typ=2) print(anova_table)
```

	sum_sq	df	F	PR(>F)
C(Program)	245.326539	2.0	61.462398	2.287256e-17
Residual	173.629808	87.0	NaN	NaN

The one-way ANOVA results show that there is a statistically significant difference in mean weight loss among the three exercise programs.

The F-statistic is 61.46, which is quite large. The p-value is 2.29×10^{-17} , which is much smaller than 0.05.

Since: $p < 0.05$

$p < 0.05$ we

reject the null

hypothesis

that all
program
means are
equal.

This indicates that at least one exercise program produces a significantly different average weight loss compared to the others.

```
In [43]: #Post Hoc Test (Tukey HSD) tukey =  
pairwise_tukeyhsd(  
endog=df['WeightLoss'],  
groups=df['Program'], alpha=0.05  
) print(tukey)
```

Multiple Comparison of Means - Tukey HSD, FWER=0.05

```
=====
```

		group1	group2	meandiff	p-adj	lower	upper	reject
	A	B	3.1005	0.0	2.2307	3.9702	True	
	A	C	-0.6985	0.1406	-1.5682	0.1713	False	
	B	C	-3.7989	0.0	-4.6687	-2.9292	True	

```
-----
```

Interpretation of Tukey HSD Test


The Tukey HSD post hoc test was conducted to identify which exercise programs differ significantly in mean weight loss.

Results show:

Program B differs significantly from Program A ($p < 0.05$). Program B also differs significantly from Program C ($p < 0.05$). There is no significant difference between Program A and Program C ($p = 0.1406$).


Program B produced the highest average weight loss, indicating that it is significantly more effective than both Programs A and C.

```
In [44]: tukey.plot_simultaneous() plt.show()
```




In [46]: In `residuals = model.resid`

[47]: `sns.histplot(residuals, kde=True) plt.title("Residual Distribution") plt.show()`



In [48]: `sm.qqplot(residuals, line='45') plt.show()`



The residual is quite normally distributed

In []:

Problem 3

Consider the Anemia Levels in Nigeria dataset: (Kaggle Dataset - Factors Affecting Children Anemia Level), which can be downloaded from Kaggle. The dataset comes from the 2018 Nigeria Demographic and Health Surveys (NDHS). It explores the impact of mothers' age and socioeconomic factors on anemia levels among children aged 0–59 months across Nigeria's 36 states and the Federal Capital Territory. Based on the above data, create a contingency table, apply chi-square test, formulate hypothesis, create a contribution diagram, and interpret the findings.

In [1]:

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import chi2_contingency
```

```
In [2]: #load dataset df = pd.read_csv("/content/children  
anemia.csv")
```

```
In [3]: df.columns
```

```
Out[3]: Index(['Age in 5-year groups', 'Type of place of residence', 'Highest educational  
level', 'Wealth index combined',  
'Births in last five years', 'Age of respondent at 1st birth',  
'Hemoglobin level adjusted for altitude and smoking (g/dl - 1 decima l)',  
'Anemia level',  
'Have mosquito bed net for sleeping (from household questionnaire)',  
'Smokes cigarettes', 'Current marital status',  
'Currently residing with husband/partner', 'When child put to breast',  
'Had fever in last two weeks',  
'Hemoglobin level adjusted for altitude (g/dl - 1 decimal)',  
'Anemia level.1', 'Taking iron pills, sprinkles or syrup'], dtype='object')
```

```
In [4]: #Select Required Variables  
data = df[[  
    'Highest educational level',  
    'Anemia level'  
]]
```

```
In [5]: #Remove Missing Values data =  
data.dropna()
```

```
In [6]: data.shape
```

```
Out[6]: (13136, 2)
```

```
[7]: #View Unique Categories  
print(data['Highest educational level'].unique()) print(data['Anemia  
level'].unique())
```

```
['Secondary' 'Higher' 'No education' 'Primary']  
['Moderate' 'Mild' 'Not anemic' 'Severe']
```

```
In [8]: #Create Contingency Table table =  
pd.crosstab( data['Highest educational level'],  
            data['Anemia level']  
) print(table)
```

```
Anemia level      Mild Moderate Not anemic Severe
```

Highest educational level				
Higher	296	238	591	14
No education	1450	1769	1874	113
Primary	608	645	900	42
Secondary	1240	1322	1972	62

In [9]:

```
#Percentage Table

row_percent = pd.crosstab( data['Highest
educational level'],
data['Anemia level'],
normalize='index'
) * 100

print(row_percent.round(2))
```

Anemia level	Mild	Moderate	Not anemic	Severe
Highest educational level				
Higher	25.99	20.90	51.89	1.23
No education	27.85	33.98	36.00	2.17
Primary	27.70	29.38	41.00	1.91
Secondary	26.98	28.76	42.91	1.35

Hypotheses **Null Hypothesis H0** :Educational level and anemia level are independent

Alternative Hypothesis H1 :Educational level and anemia level are associated

In [10]:

```
#Chi-Square Test
chi2, p, dof, expected = chi2_contingency(table)

print("Chi-square Statistic:", chi2) print("Degrees of
Freedom:", dof)
print("P-value:", p)
```

Chi-square Statistic: 142.86240492075743

Degrees of Freedom: 9

P-value: 2.644147441640121e-26

Interpretation of Chi-Square Test

The chi-square test results indicate a statistically significant association between educational level and anemia level.

Chi-square statistic = 142.86 Degrees of freedom = 9 p-value = 2.64×10^{-26}

Since: $p < 0.05$

we reject the null hypothesis of independence.

This suggests that anemia levels vary significantly across different educational levels, indicating that educational status may be an important factor related to child anemia outcomes.

In [11]:

```
#Chi-Square Contribution Matrix
contribution = ((table - expected)**2) / expected
print(contribution.round(2))
```

Anemia level	Mild	Moderate	Not anemic	Severe
Highest educational level				
Higher	0.78	32.96	35.54	1.82
No education	0.46	23.91	27.49	5.03
Primary	0.09	0.55	0.08	0.30
Secondary	0.24	3.37	5.87	4.38

Interpretation of Contribution Matrix

The contribution matrix shows which combinations of educational level and anemia level contribute most to the chi-square statistic.

The largest contributions are observed for:

Higher education × Not anemic (35.54) Higher education × Moderate anemia (32.96) No education × Not anemic (27.49) No education × Moderate anemia (23.91)

These large values indicate that these categories are mainly responsible for the significant association between educational level and anemia status.


The results suggest that children's anemia outcomes differ considerably across

maternal educational groups, especially between higher education and no education categories.

In [12]:

```
#Contribution Heatmap/diagram plt.figure(figsize=(10,6))

sns.heatmap( contribution,
              annot=True,
              cmap='Reds'
) plt.title("Chi-Square Contribution Heatmap")
plt.show()
```



Interpretation of Heatmap

The heatmap visualizes the contribution of each cell to the overall chi-square statistic.

Darker red cells indicate higher contributions and stronger deviations from expected frequencies. The cells for: Higher education × Not anemic Higher education × Moderate anemia No education × Not anemic No education × Moderate anemia show the largest contributions.

This means these categories are the main drivers of the significant association between educational level and anemia status.

In contrast, the Primary education group shows very small contribution values, indicating little influence on the chi-square result.

In []:

Problem 4

Suppose there are three drug treatments (Drug A, Drug B, and Drug C) with the outcome of disease or no disease. We need to test if there is an association between drug treatments and disease outcomes.

Treatment No Disease Disease Drug A 40 10 Drug B 10 40 Drug C 25 25

Perform Fisher's exact test for the above contingency table, conduct a post-hoc test for Fisher's exact test, and interpret your findings.

In [1]:

```
import pandas as pd
import numpy as np

from scipy.stats import fisher_exact
from scipy.stats import chi2_contingency
from statsmodels.stats.multitest import multipletests

import seaborn as sns
import matplotlib.pyplot as plt
```

In [2]:

```
#Create Contingency Table table =
pd.DataFrame({
    'No_Disease': [40, 10, 25],
    'Disease': [10, 40, 25]
},
index=['Drug_A', 'Drug_B', 'Drug_C'])
print(table)
```

	No_Disease	Disease
Drug_A	40	10
Drug_B	10	40
Drug_C	25	25

In [3]:

```
#Overall Association Test
#Since Fisher's Exact Test in scipy works directly only for 2x2 tables, we fir chi2, p, dof,
expected = chi2_contingency(table)

print("Chi-square Statistic:", chi2)
print("Degrees of Freedom:", dof)
print("P-value:", p)
```

Chi-square Statistic: 36.0

Degrees of Freedom: 2

P-value: 1.522997974471264e-08

Interpretation of Overall Association Test

The chi-square test results indicate a statistically significant association between drug treatment and disease outcome.

Chi-square statistic = 36.0 Degrees of freedom = 2 p-value = 1.52×10^{-8}

Since: $p < 0.05$

we reject the null hypothesis of independence.

This suggests that disease outcomes differ significantly across the three drug treatments, indicating that treatment type may influence disease status.

Pairwise Fisher's Exact Tests

We compare drugs pairwise:

A vs B A vs C B vs C

Hypotheses

Null Hypothesis

H_0 :Treatment and disease outcome are independent

Alternative Hypothesis

H_1 :Treatment and disease outcome are associated

In [4]:

```
#A vs B
AB = np.array([
    [40, 10],
    [10, 40]
]) oddsratio_AB, p_AB = fisher_exact(AB)

print("A vs B")
print("Odds Ratio:", oddsratio_AB) print("P-value:", p_AB)
```

A vs B

Odds Ratio: 16.0

P-value: 2.2221008148312167e-09

In [5]:

```
#A vs C
AC = np.array([
    [40, 10],
    [25, 25] ])

oddsratio_AC, p_AC = fisher_exact(AC)

print("A vs C") print("Odds Ratio:", oddsratio_AC)
print("P-value:", p_AC)
```

A vs C

Odds Ratio: 4.0

P-value: 0.0030520593364333462

In [6]:

```
#B vs C
BC = np.array([
    [10, 40],
    [25, 25]
]) oddsratio_BC, p_BC = fisher_exact(BC)

print("B vs C")
print("Odds Ratio:", oddsratio_BC) print("P-value:", p_BC)
```

B vs C

Odds Ratio: 0.25

P-value: 0.0030520593364333462

In [8]:

```
#Multiple Testing Correction #Bonferroni

correction:

p_values = [p_AB, p_AC, p_BC]

adjusted = multipletests(p_values,
    method='bonferroni'
)

print("Adjusted P-values:")
print(adjusted[1])
```

Adjusted P-values:

[6.66630244e-09 9.15617801e-03 9.15617801e-03]

In [9]:

```
#Create Summary Table results =  
pd.DataFrame({  
    'Comparison': ['A vs B', 'A vs C', 'B vs C'],  
    'P-value': p_values,  
    'Adjusted P-value': adjusted[1]  
}) print(results)
```

	Comparison	P-value	Adjusted P-value
0	A vs B	2.222101e-09	6.666302e-09
1	A vs C	3.052059e-03	9.156178e-03
2	B vs C	3.052059e-03	9.156178e-03

Interpretation of Pairwise Fisher's Exact Tests

The pairwise Fisher's exact tests with Bonferroni correction show significant differences between all drug treatment pairs.


Drug A vs Drug B: adjusted p-value = 6.67×10^{-9} Drug A vs Drug C: adjusted pvalue = 0.009 Drug B vs Drug C: adjusted p-value = 0.009

Since all adjusted p-values are less than 0.05, we conclude that the disease outcomes differ significantly among all treatment groups.

The strongest difference is observed between Drug A and Drug B, suggesting these two treatments have substantially different effects on disease outcomes.

In [10]:

```
plt.figure(figsize=(6,4))  
  
sns.heatmap( table,  
    annot=True,  
    cmap='Blues',  
    fmt='d'  
) plt.title("Drug Treatment vs Disease Outcome")  
plt.show()
```



In []:

Problem 5

A researcher is interested in how variables, such as GRE (Graduate Record Exam scores), GPA (grade point average), and prestige of the undergraduate institution, affect admission into graduate school. The response variable, admit/don't admit, is a binary variable. The dataset is taken from:

UCLA Binary Dataset

Fit logistic generalized linear models (GLMs) to identify the factors affecting admission into graduate school. Interpret the results.

In [1]:

```
import pandas as pd
import numpy as np

import statsmodels.api as sm
import statsmodels.formula.api as smf

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.metrics import (confusion_matrix,
                             classification_report,
                             roc_curve, roc_auc_score)
```

In [2]:

```
#load data
url = "https://stats.idre.ucla.edu/stat/data/binary.csv"
df = pd.read_csv(url)
```

In [3]:

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 400 entries, 0 to 399 Data

columns (total 4 columns):

Column Non-Null Count Dtype

--- ---

0 admit 400 non-null int64


1 gre 400 non-null int64

2 gpa 400 non-null float64


3 rank 400 non-null int64 dtypes: float64(1), int64(3) memory usage: 12.6 KB

```
In [4]: ##Exploratory Data Analysis df.describe()
sns.boxplot(x='admit', y='gre', data=df)


plt.title("GRE vs Admission") plt.show()
```



```
In [5]: sns.boxplot(x='admit', y='gpa', data=df)
plt.title("GPA vs Admission") plt.show()
```



In [6]: `sns.countplot(x='rank', hue='admit', data=df)`
`plt.title("Admission by Rank") plt.show()`



In [7]: `#Convert Rank into Categorical Variable df['rank'] =
df['rank'].astype('category')`

Fit Logistic Regression Model Model:

`admit~gre+gpa+rank admit~gre+gpa+rank`

In [8]: `model = smf.logit(
 'admit ~ gre + gpa + rank',
 data=df
).fit()`

Optimization terminated successfully. Current
function value: 0.573147
Iterations 6

In [9]: `print(model.summary())`

Logit Regression Results

```
=====
=====
Dep. Variable:          admit  No. Observations:          400
Model:                Logit  Df Residuals:              394
```

Method:	MLE	Df Model:	5
Date:	Sun, 10 May 2026	Pseudo R-squ.:	0.08292
Log-Likelihood:	-229.26	converged:	True
Covariance Type:	nonrobust	LLR p-value:	7.578e-08

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-3.9900	1.140	-			
rank[T.2]	-0.6754	0.316	-2.134	0.033	-1.296	
rank[T.3]	-1.3402	0.345	-3.881	0.000	-2.017	-0.663
rank[T.4]	-1.5515					
gre	0.0023	0.001	2.070	0.038		
gpa	0.8040	0.332	2.423	0.015	0.154	1.454

Interpretation of Logistic Regression Results

The logistic regression model examines how GRE score, GPA, and undergraduate institution rank affect the probability of admission into graduate school.

The overall model is statistically significant:

Likelihood Ratio p-value = 7.58×10^{-8}

This indicates that the predictors collectively help explain admission outcomes.

Interpretation of Predictors

GRE Score

Coefficient = 0.0023 p-value = 0.038

Since $p < 0.05$, GRE is a statistically significant predictor.

A higher GRE score increases the likelihood of admission.

GPA

Coefficient = 0.8040 p-value = 0.015

GPA is also statistically significant.

Students with higher GPAs are more likely to be admitted into graduate school.

Institution Rank

Compared to rank 1 institutions (reference category):

Rank 2 coefficient = -0.6754 Rank 3 coefficient = -1.3402 Rank 4 coefficient = -1.5515

All coefficients are negative and statistically significant.

This suggests that students from lower prestige institutions have lower odds of admission compared to students from rank 1 institutions.

Model Performance

Pseudo $R^2 = 0.0829$

This indicates the model explains a modest proportion of variability in admission outcomes, which is common in logistic regression for real-world admission data.

Overall Conclusion

The analysis suggests that: higher GRE scores, higher GPA, and graduating from more prestigious institutions increase the probability of admission into graduate school.

In [10]:

```
#odd ratios
odds_ratios = np.exp(model.params) print(odds_ratios)
```

```
Intercept    0.018500
rank[T.2]    0.508931
rank[T.3]    0.261792
rank[T.4]    0.211938 gre
1.002267 gpa    2.234545
dtype: float64
```

Interpretation of Odds Ratios

The odds ratios explain how each predictor affects the odds of admission to graduate school.

GRE

Odds Ratio = 1.002

This means that for each one-unit increase in GRE score, the odds of admission increase by about 0.2%, holding other variables constant.

GPA

Odds Ratio = 2.23

This indicates that a one-unit increase in GPA multiplies the odds of admission by about 2.23 times.

Higher GPA substantially increases admission chances.

Institution Rank

Rank 1 is the reference category.

Compared to rank 1 institutions:

Rank 2 OR = 0.51 Rank 3 OR = 0.26 Rank 4 OR = 0.21

Since these odds ratios are less than 1, students from lower-ranked institutions have lower odds of admission.

For example:

Students from rank 4 institutions have about 79% lower odds of admission compared to students from rank 1 institutions.

In []:

Problem 6

The number of awards earned by students at one high school. Predictors of the number of awards earned include the type of program in which the student was enrolled (e.g., vocational, general, or academic) and the score on their final exam in math. The dataset is taken from:

UCLA Poisson Dataset

Fit the Poisson generalized linear model (GLM) to identify the factors associated with the number of awards earned by students at one high school. Interpret the results.

In [1]:

```
import pandas as pd
import numpy as np

import statsmodels.api as sm
import statsmodels.formula.api as smf

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
```

```
In [2]: #load dataset
url = "https://stats.idre.ucla.edu/stat/data/poisson_sim.csv" df =
pd.read_csv(url)
```

```
In [3]: df.head() df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199 Data
columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id          200 non-null    int64
1   num_awards  200 non-null    int64
2   prog        200 non-null    int64
3   math        200 non-null    int64 dtypes: int64(4) memory usage: 6.4 KB
```

```
In [4]: #Convert Program to Categorical df['prog'] =
df['prog'].astype('category')
```


```
In [5]: #Exploratory Data Analysis df.describe()
```

Out[5]:

	id	num_awards	math
count	200.000000	200.000000	200.000000
mean	100.500000	0.630000	52.645000
std	57.879185	1.052921	9.368448
min	1.000000	0.000000	33.000000
25%	50.750000	0.000000	45.000000
50%	100.500000	0.000000	52.000000
75%	150.250000	1.000000	59.000000
max	200.000000	6.000000	75.000000


In [6]:

```
#Distribution of Awards  
sns.countplot(x='num_awards', data=df) plt.title("Distribution of  
Number of Awards") plt.show()
```




In [7]:

```
#Awards by Program sns.boxplot(  
  
    x='prog', y='num_awards',  
    data=df  
) plt.title("Awards by Program")  
plt.show()
```



In [8]:

```
#Math Score vs Awards  
sns.scatterplot( x='math',  
y='num_awards',  
hue='prog', data=df  
) plt.title("Math Score vs Awards")  
plt.show()
```



```
In [9]: In #Fit Poisson Regression Model
model = smf.glm( formula='num_awards ~
prog + math',
data=df,
family=sm.families.Poisson() ).fit()

print(model.summary())
```

[10]:

Generalized Linear Model Regression Results

```
=====
Dep. Variable: num_awards No. Observations: 200
Model: GLM Df Residuals: 196
Model Family: Poisson Df Model: 3
```


Coefficient = 0.3698 p-value = 0.402

Program 3 is not statistically significant because $p > 0.05$.

There is no strong evidence that Program 3 differs from Program 1 in the expected number of awards.

Interpretation of Coefficients in Poisson Regression

Since Poisson regression uses a log link:

$$\log(\mu) = \beta_0$$

$$+ \beta_1$$

$$X_1$$

$$+ \dots$$

$$\log(\mu) = \beta_0$$

$$+ \beta_1$$

$$X_1$$

$$+ \dots$$

Exponentiating coefficients gives Incident Rate Ratios (IRR).

In []:

Problem 7

School administrators study the attendance behavior of high-school juniors at two schools. Predictors of the number of days of absence include the type of program in which the student is enrolled and a standardized math test score. The dataset is taken from:

UCLA Negative Binomial Dataset

Fit the negative binomial generalized linear model (GLM) to identify the factors associated with the number of days absent by students. Interpret the results.


```
In [1]: import pandas as pd
import numpy as np

import statsmodels.api as sm
import statsmodels.formula.api as smf

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
```

```
In [2]: url = "https://stats.idre.ucla.edu/stat/stata/dae/nb_data.dta"
df = pd.read_stata(url)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	id	gender	math	daysabs	prog
0	1001.0	male	63.0	4.0	2.0

```
df.info()
```

1	1002.0	male	27.0	4.0	2.0
2	1003.0	female	20.0	2.0	2.0
3	1004.0	female	16.0	3.0	2.0
4	1005.0	female	2.0	3.0	2.0

```
In [4]:
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 314 entries, 0 to 313
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0   id      314 non-null     float32
1   gender  314 non-null     category
2   math    314 non-null     float32
3   daysabs 314 non-null     float32
```

4 prog 314 non-null float32 dtypes: category(1), float32(4) memory usage: 5.5 KB

```
In [5]: #Convert Program to Category
df['prog'] = df['prog'].astype('category')
```

```
In [6]: #Exploratory Data Analysis
#Summary Statistics df.describe()
```

```
Out[6]:
```

	id	math	daysabs
count	314.000000	314.000000	314.000000

```
#Check Overdispersion print(
    "Mean:",
    df['daysabs'].mean()
)

print(
    "Variance:", df['daysabs'].var()
)
```

mean	1575.910889	48.267517	5.955414
std	502.314789	25.362385	7.036952
min	1001.000000	1.000000	0.000000
25%	1079.250000	28.000000	1.000000
50%	1158.500000	48.000000	4.000000
75%	2077.750000	70.000000	8.000000
max	2157.000000	99.000000	35.000000

```
In [7]:
```

Mean: 5.955414

Variance: 49.518699645996094 since Variance >

Mean so overdispersion exists.

This justifies using Negative Binomial regression instead of Poisson regression.

In [8]:

```
#Fit Negative Binomial Model

#Model:daysabs~prog+math model =
smf.glm( formula='daysabs ~ prog + math',
        data=df,
        family=sm.families.NegativeBinomial() ).fit()
```

```
/usr/local/lib/python3.12/dist-packages/statsmodels/genmod/families/family.py:1
367: ValueWarning: Negative binomial dispersion parameter alpha not set. Using default value
alpha=1.0. warnings.warn("Negative binomial dispersion parameter alpha not ")
```

In [9]:

```
print(model.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:      daysabs  No. Observations:      314
Model:              GLM    Df Residuals:           310
Model Family:      NegativeBinomial  Df Model:         3
Link Function:      Log    Scale:                1.0000
Method:            IRLS   Log-Likelihood:         -865.68
Date:              Sun, 10 May 2026  Deviance:         350.98 Time:           12:03:50
Pearson chi2:       331.
No. Iterations:     6   Pseudo R-squ. (CS):         0.1926
Covariance Type:    nonrobust
=====
```

```
===== coef  std err      z  P>|z|  [0.025  0.975]
-----Intercept  2.6150    0.200
13.054  0.000    2.222   3.008 prog[T.2.0] -0.4408    0.185  -2.379    0.017  -0.804
-0.078 prog[T.3.0] -1.2786    0.203  -6.284    0.000  -1.677  -0.880 math  -0.0060
0.003          -2.358          0.018          -0.011          -0.001
=====
```

Interpretation of Negative Binomial Regression Results

The negative binomial regression model examines how program type and math score affect the number of days absent by students.

The overall model appears meaningful, with a pseudo R^2 of 0.193, indicating that the predictors explain part of the variation in absenteeism.

Interpretation of Predictors

Math Score

Coefficient = -0.0060 p-value = 0.018 Math

score is statistically significant.

The negative coefficient indicates that students with higher math scores tend to have fewer days absent.

Program Type

Reference category = Program 1

Program 2

Coefficient = -0.4408 p-value = 0.017

Program 2 is statistically significant.

Students in Program 2 are expected to have fewer absences compared to students in Program 1.

Program 3

Coefficient = -1.2786 p-value < 0.001

Program 3 is highly significant.

Students in Program 3 are expected to have substantially fewer absences compared to students in Program 1.

Interpretation of Coefficients

Negative Binomial regression uses a log link:

$$\log(\mu) = \beta_0 + \beta_1 X_1 + \dots$$

Negative coefficients reduce the expected count of absences.

Overall Conclusion

The analysis suggests that:

higher math scores are associated with lower absenteeism, and students enrolled in Programs 2 and 3 tend to have significantly fewer absences than students in the reference program.

In []:

Problem 8

The state wildlife biologists want to model how many fish are being caught by fishermen at a state park. Visitors are asked how long they stayed, how many people were in the group, whether there were children in the group, and how many fish were caught. Some visitors do not fish, but there is no data on whether a person fished or not. Some visitors who did fish did not catch any fish, so there are excess zeros in the data because of the people that did not fish. The dataset is taken from:

UCLA Fish Dataset

Fit the Zero-Inflated Poisson (ZIP) regression generalized linear model (GLM) to identify the factors associated with the number of fish caught. Interpret the results.

In [1]:

```
import pandas as pd
import numpy as np

import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.discrete.count_model import ZeroInflatedPoisson

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
```

In [2]:

```
url = "https://stats.idre.ucla.edu/stat/data/fish.csv"
df = pd.read_csv(url)
```

In [3]:

```
df.head()
```

Out[3]:

	nofish livebait camper persons child					xb		zg count
0	1	0	0	1	0	-0.896315	3.050405	0
1	0	1	1	1	0	-0.558345	1.746149	0
2	0	1	0	1	0	-0.401731	0.279939	0
3	0	1	1	2	1	-0.956298	-0.601526	0

4 0 1 0 1 0 0.436891 0.527709 1

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249 Data
columns (total 8 columns):
#   Column   Non-Null Count  Dtype
---  -
0  nofish   250 non-null     int64
1  livebait 250 non-null     int64
2  camper   250 non-null     int64
3  persons  250 non-null     int64
4  child    250 non-null     int64
5  xb       250 non-null     float64
6  zg       250 non-null     float64
7  count    250 non-null     int64 dtypes: float64(2), int64(6) memory usage: 15.8 KB
```

In [5]: `df.describe()`

Out[5]:

	nofish	livebait	camper	persons	child	xb
count	250.000000	250.000000	250.000000	250.000000	250.000000	250.000000
mean	0.296000	0.864000	0.588000	2.52800	0.684000	0.973796
std	0.457407	0.343476	0.493182	1.11273	0.850315	1.440277
min	0.000000	0.000000	0.000000	1.00000	0.000000	-3.275050
25%	0.000000	1.000000	0.000000	2.00000	0.000000	0.008267
50%	0.000000	1.000000	1.000000	2.00000	0.000000	0.954550
75%	1.000000	1.000000	1.000000	4.00000	1.000000	1.963855
max	1.000000	1.000000	1.000000	4.00000	3.000000	5.352674

```
In [6]: #Check Excess Zeros
zeros = (df['count'] == 0).sum() print("Number of
Zeros:", zeros)

print(
    "Percentage of Zeros:",
    zeros / len(df) * 100 )
```

Number of Zeros: 142

Percentage of Zeros: 56.8

```
In [7]: #Prepare Variables #Response
```

```
In [8]: variable:
```

```
y = df['count']
```

```
In [9]: In #Predictor variables:
```

```
X = df[['child', 'persons', 'camper']]
```

```
#Add constant:
```

```
[10]: X = sm.add_constant(X)
```

```
#Fit Zero-Inflated Poisson Model #Model:
```

```
count~child+persons+camper
```

```
zip_model = ZeroInflatedPoisson(
    endog=y, exog=X, exog_infl=X,
    inflation='logit'
).fit()
```

Optimization terminated successfully.

Current function value: 3.010926

Iterations: 29

Function evaluations: 31

Gradient evaluations: 31

```
In [11]: print(zip_model.summary())
```

ZeroInflatedPoisson Regression Results

```

=====
Dep. Variable:          count  No. Observations:          250
Model:          ZeroInflatedPoisson  Df Residuals:          246
Method:          MLE  Df Model:          3
Date:          Sun, 10 May 2026  Pseudo R-squ.:          0.3321  Time:          12:14:42
Log-Likelihood:          -752.73  converged:          True  LL-Null:          -1127.0
Covariance Type:          nonrobust  LLR p-value:          6.123e-162
=====

```

```

=====
==== coef  std err      z    P>|z|    [0.025
0.975] -----
----
inflate_const    1.6636    0.503    3.308    0.001    0.678
2.649
inflate_child    1.9046    0.322    5.915    0.000    1.273
2.536
inflate_persons  -0.9228    0.197   -4.689    0.000   -1.308
-0.537
inflate_camper   -0.8336    0.350   -2.385    0.017   -1.519
-0.148
const           -0.7983    0.165   -4.830    0.000   -1.122
-0.474
child           -1.1367    0.091  -12.464    0.000   -1.315
-0.958
persons          0.8290    0.043   19.415    0.000    0.745
0.913
camper           0.7243    0.092    7.854    0.000    0.544
0.905
=====
=====
=====

```

Interpretation of Zero-Inflated Poisson (ZIP) Regression Results

The Zero-Inflated Poisson model was used because the dataset contains excess zeros, meaning many visitors caught no fish, possibly because some individuals did not fish at all.

The overall model is highly significant:

Likelihood Ratio p-value = 6.12×10^{-162} Pseudo $R^2 = 0.332$

This indicates that the predictors explain a substantial portion of the variation in fish counts.

Count Model Interpretation

These coefficients explain factors affecting the number of fish caught.

Number of Children (child) Coefficient = -1.1367 p-value < 0.001

This variable is statistically significant.

Groups with more children are expected to catch fewer fish. Number of

Persons (persons) Coefficient = 0.8290 p-value < 0.001

This variable is statistically significant.

Larger groups are associated with higher expected fish counts.

Camper Status (camper)

Coefficient = 0.7243 p-value < 0.001

Campers are expected to catch more fish compared to non-campers.

Inflation Model Interpretation

The inflation part models the probability of excess zeros. inflate_child

Positive coefficient:

More children increase the probability of belonging to the “always zero” group.

Meaning: groups with more children are more likely not to fish at all. inflate_persons

Negative coefficient:

Larger groups are less likely to belong to the excess-zero group.

Meaning: larger groups are more likely to participate in fishing. inflate_camper

Negative coefficient:

Campers are less likely to belong to the excess-zero group.

Meaning: campers are more likely to actually fish.

Overall Conclusion

The ZIP regression model indicates that:

larger groups and campers tend to catch more fish, groups with more children tend to catch fewer fish, and the excess zeros are partly explained by visitors who likely did not participate in fishing activities.

In []:

Problem 9

The state wildlife biologists want to model how many fish are being caught by fishermen at a state park. Visitors are asked how long they stayed, how many people were in the group, whether there were children in the group, and how many fish were caught. Some visitors do not fish, but there is no data on whether a person fished or not. Some visitors who did fish did not catch any fish, so there are excess zeros in the data because of the people that did not fish. The dataset is taken from:

UCLA Fish Dataset

Fit the Zero-Inflated Negative Binomial (ZINB) regression generalized linear model (GLM) to identify the factors associated with the number of fish caught. Interpret the results.

```
In [1]: import pandas as pd
import numpy as np
import statsmodels.api as sm
from statsmodels.discrete.count_model import ZeroInflatedNegativeBinomialP

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
```

```
In [17]: url = "https://stats.idre.ucla.edu/stat/data/fish.csv"
df = pd.read_csv(url)
```

```
In [18]: df.head()
```

```
Out[18]:
```

	nofish	livebait	camper	persons	child	xb	zg	count
0	1	0	0	1	0	-0.896315	3.050405	0
1	0	1	1	1	0	-0.558345	1.746149	0
2	0	1	0	1	0	-0.401731	0.279939	0
3	0	1	1	2	1	-0.956298	-0.601526	0
4	0	1	0	1	0	0.436891	0.527709	1

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 250 entries, 0 to 249 Data
columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	nofish	250 non-null	int64
1	livebait	250 non-null	int64
2	camper	250 non-null	int64
3	persons	250 non-null	int64
4	child	250 non-null	int64
5	xb	250 non-null	float64
6	zg	250 non-null	float64
7	count	250 non-null	int64

dtypes: float64(2), int64(6) memory usage: 15.8 KB

In [20]: df.describe()

	nofish	livebait	camper	persons	child	xb
count	250.000000	250.000000	250.000000	250.000000	250.000000	250.000000
mean	0.296000	0.864000	0.588000	2.528000	0.684000	0.973796
std	0.457407	0.343476	0.493182	1.11273	0.850315	1.440277
min	0.000000	0.000000	0.000000	1.000000	0.000000	-3.275050
25%	0.000000	1.000000	0.000000	2.000000	0.000000	0.008267

```
#Check Excess Zeros
zeros = (df['count'] == 0).sum() print("Number of  
Zeros:", zeros)

print(
    "Percentage of Zeros:",
    zeros / len(df) * 100 )
```

50%	0.000000	1.000000	1.000000	2.000000	0.000000	0.954550
75%	1.000000	1.000000	1.000000	4.000000	1.000000	1.963855
max	1.000000	1.000000	1.000000	4.000000	3.000000	5.352674

In [21]:

Number of Zeros: 142
Percentage of Zeros: 56.8

```
In [22]: #Mean and Variance print(
          "Mean:",
          df['count'].mean()
        )

        print(
          "Variance:", df['count'].var()
        )
```

Mean: 3.296

Variance: 135.37387951807236 since Variance >>

Mean so overdispersion exists This supports using
ZINB instead of ZIP.

```
In [23]: #Prepare Variables #Response

          variable:

          y = df['count']
```

```
In [24]: #Predictors:

          X = df[['child', 'persons', 'camper']]
```

```
In [25]: X = sm.add_constant(X)
```

```
In [15]: #Fit ZINB Model

          #Model:count~child+persons+camper zinb_model
          = ZeroInflatedNegativeBinomialP( endog=y,
          exog=X, exog_infl=X, inflation='logit'
          ).fit()
```

Current function value: 1.620380

Iterations: 35

Function evaluations: 37

Gradient evaluations: 37

```
print(zinb_model.summary())
```

```
/usr/local/lib/python3.12/dist-packages/statsmodels/discrete/discrete_model.p
```

```
y:3937: RuntimeWarning: invalid value encountered in log a1 *
```

```
np.log(a1) + y * np.log(mu) -
```

```
/usr/local/lib/python3.12/dist-packages/statsmodels/discrete/discrete_model.p
```

```
y:3938: RuntimeWarning: invalid value encountered in log
```

```
(y + a1) * np.log(a2))
```

```
/usr/local/lib/python3.12/dist-packages/scipy/optimize/_optimize.py:1330: Optim izeWarning:
```

```
Maximum number of iterations has been exceeded. res = _minimize_bfgs(f, x0, args, fprime,  
callback=callback, **opts)
```

```
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:595: HessianI
```

```
nversionWarning: Inverting hessian failed, no bse or cov_params available
```

```
warnings.warn('Inverting hessian failed, no bse or cov_params '
```

```
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: Converge nceWarning:
```

```
Maximum Likelihood optimization failed to converge. Check mle_retva
```

```
ls warnings.warn("Maximum Likelihood optimization failed to "
```

In [26]:

ZeroInflatedNegativeBinomialP Regression Results

```

=====
===== Dep. Variable:          count  No. Observations:
250 Model:          ZeroInflatedNegativeBinomialP  Df Residuals:
246 Method:          MLE  Df Model:
3 Date:              Sun, 10 May 2026  Pseudo R-squ.:
0.1278 Time:         12:40:17  Log-Likelihood:
-405.09 converged:   False  LL-Null: -464.44 Covariance
Type:               nonrobust  LLR p-value:
1.479e-25
=====

```

```

=====
===== coef  std err      z    P>|z|    [0.025
0.975] -----
----
inflate_const  -0.4892      nan      nan      nan      nan
nan
inflate_child   0.0378      nan      nan      nan      nan
nan
inflate_persons -1.5033      nan      nan      nan      nan
nan
inflate_camper  -0.4691      nan      nan      nan      nan
nan
const          -1.5122      nan      nan      nan      nan
nan
child          -1.7885      nan      nan      nan      nan
nan
persons        1.0343      nan      nan      nan      nan
nan
camper         0.6020      nan      nan      nan      nan
nan
alpha          2.0825      nan      nan      nan      nan
nan
=====
=====
=====

```

Interpretation of Zero-Inflated Negative Binomial (ZINB) Results

The Zero-Inflated Negative Binomial model was fitted to analyze fish catch counts while accounting for:

excess zeros overdispersion in the data

The overall model is statistically significant:

Likelihood Ratio p-value = 1.48×10^{-25}

This suggests that the predictors collectively help explain fish catch counts.

In []:

Problem 10

A study of length of hospital stay, in days, as a function of age, kind of health insurance, and whether or not the patient died while in the hospital. Length of hospital stay is recorded as a minimum of at least one day. The dataset is taken from:

UCLA Zero-Truncated Poisson Dataset

Fit the Zero-Truncated Poisson regression generalized linear model (GLM) to identify the factors associated with hospital length of stay. Interpret the results.

```
In [1]: import pandas as pd import numpy as np import statsmodels.api as sm from
statsmodels.discrete.truncated_model import TruncatedLFPoisson

import seaborn as sns import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
```

```
In [5]: pip install pyreadstat
```

Collecting pyreadstat

Downloading pyreadstat-1.3.4-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_28_x86_64.whl.metadata (1.2 kB)

Requirement already satisfied: narwhals>=2.10.1 in /usr/local/lib/python3.12/dist-packages (from pyreadstat) (2.20.0)

Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from pyreadstat) (2.0.2)

Downloading pyreadstat-1.3.4-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_28_x86_64.whl (2.6 MB)

2.6/2.6 MB 27.6 MB/s eta 0:00:00

Installing collected packages: pyreadstat

Successfully installed pyreadstat-1.3.4

```
In [7]: import pandas as pd
import pyreadstat import
urllib.request
```

```
In [8]: url = "https://stats.idre.ucla.edu/stat/data/ztp.dta"

urlib.request.urlretrieve( url,
    "ztp.dta"
)
```

```
Out[8]:
('ztp.dta', df, meta = pyreadstat.read_dta( "ztp.dta"
)
```

<http.client.HTTPMessage at 0x7b0e8dee7080>)

```
In [9]:
```

```
In [10]: df.head()
```

```
Out[10]:      stay age hmo died
```

0	4	4	0	0.0
1	9	4	1	0.0
2	3	7	1	1.0
3	9	6	0	0.0
4	1	7	0	1.0

```
In [11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1493 entries, 0 to 1492 Data
columns (total 4 columns):
#   Column Non-Null Count  Dtype
---  -
0  stay    1493 non-null    int64
1  age     1493 non-null    int64
2  hmo     1493 non-null    int64
3  died    1493 non-null    float64 dtypes: float64(1), int64(3) memory usage: 46.8 KB
```

```
In [12]: #Convert Variables
df['hmo'] = df['hmo'].astype('category') df['died'] =
df['died'].astype('category')
```

In [13]: `df.describe()`

Out[13]:

	stay	age
count	1493.000000	1493.000000
mean	9.728734	5.233758
std	8.132908	1.669273
min	1.000000	1.000000
25%	4.000000	4.000000

```
#Check Zero Counts print(  
    "Number of zeros:",  
    (df['stay'] == 0).sum() )
```

50%	8.000000	5.000000
75%	13.000000	6.000000
max	74.000000	9.000000

In [14]:

Number of zeros: 0

In [15]:

```
#Prepare Variables #Response
```

```
variable:
```

```
y = df['stay']
```

In [19]:

```
X = pd.get_dummies( df[['age', 'hmo',  
    'died']],  
    drop_first=True )
```

```
X = X.astype(float)
```

```
X = sm.add_constant(X)
```

In [17]:

In [20]:

```
#Fit Zero-Truncated Poisson Model

#Model:stay~age+hmo+died
ztp_model = TruncatedLFPoisson(
endog=y, exog=X
).fit()
```

Optimization terminated successfully. Current
function value: 4.627461
Iterations: 4
Function evaluations: 6
Gradient evaluations: 6

In [21]:

```
print(ztp_model.summary())
```

TruncatedLFPoisson Regression Results

```
=====
Dep. Variable:          stay  No. Observations:          1493
Model:          TruncatedLFPoisson  Df Residuals:          1489
Method:          MLE  Df Model:          3
Date:          Sun, 10 May 2026  Pseudo R-squ.:          0.01294  Time:          12:54:21
Log-Likelihood:          -6908.8 converged:          True  LL-Null:          -
6999.4 Covariance Type:          nonrobust  LLR p-value:          5.022e-39
=====
```

```
===== coef  std err      z  P>|z|  [0.025  0.975]
-----
89.118  0.000   2.382   2.489 age      -0.0144   0.005  -2.868   0.004  -0.024
-0.005 hmo_1  -0.1359   0.024  -5.724   0.000  -0.182  -0.089 died_1.0  -0.2038
0.018      -11.091      0.000      -0.240      -0.168
=====
```

Interpretation of Zero-Truncated Poisson Regression Results

The Zero-Truncated Poisson regression model was used to identify factors associated with hospital length of stay, where zero-day stays are excluded from the dataset.

The overall model is statistically significant:

Likelihood Ratio p-value = 5.02×10^{-39}

This indicates that the predictors collectively help explain variation in hospital stay duration.

Interpretation of Predictors

Age:

Coefficient = -0.0144 p-value = 0.004

Age is statistically significant.

The negative coefficient suggests that higher age is associated with a slightly shorter expected hospital stay.

Insurance Type (hmo_1)

Coefficient = -0.1359 p-value < 0.001

Insurance type is statistically significant.

Patients with HMO insurance tend to have shorter hospital stays compared to the reference group.

Mortality Status (died_1.0)

Coefficient = -0.2038 p-value < 0.001

Mortality status is statistically significant.

Patients in the died = 1 category are associated with shorter expected hospital stays compared to the reference category.

Interpretation of Coefficients

The Zero-Truncated Poisson model uses a log link:

$$\log(\mu) = \beta_0 + \beta_1 X_1 + \dots$$

Negative coefficients reduce the expected hospital stay duration.

Overall Conclusion

The analysis suggests that: age, insurance

type, and mortality status

are significant predictors of hospital length of stay. The Zero-Truncated Poisson model was appropriate because the dataset excludes zero-day hospital stays.

In []:

Problem 11

A study of length of hospital stay, in days, as a function of age, kind of health insurance, and whether or not the patient died while in the hospital. Length of hospital stay is recorded as a minimum of at least one day. The dataset is taken from:

UCLA Zero-Truncated Dataset

Fit the Zero-Truncated Negative Binomial regression generalized linear model (GLM) to identify the factors associated with hospital length of stay. Interpret the results.

In [2]:

```
import pandas as pd
import numpy as np
import statsmodels.api as sm
from statsmodels.discrete.truncated_model import TruncatedLNegativeBinomialP

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error

import urllib.request
```

In [3]:

```
!pip install pyreadstat
```

Collecting pyreadstat

Downloading pyreadstat-1.3.4-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_28_x86_64.whl.metadata (1.2 kB)

Requirement already satisfied: narwhals>=2.10.1 in /usr/local/lib/python3.12/dist-packages (from pyreadstat) (2.20.0)

Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from pyreadstat) (2.0.2)

Downloading pyreadstat-1.3.4-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_28_x86_64.whl (2.6 MB)

2.6/2.6 MB 25.4 MB/s eta 0:00:00

Installing collected packages: pyreadstat

Successfully installed pyreadstat-1.3.4

In [4]:

```
import pyreadstat
```

In [5]: `url = "https://stats.idre.ucla.edu/stat/data/ztp.dta"`

Out[5]: `urllib.request.urlretrieve(url,
"ztp.dta"
('ztp.dta',)`

```
df, meta = pyreadstat.read_dta( "ztp.dta"  
)
```

<http.client.HTTPMessage at 0x7e15ad235520>) In [6]:

In [7]: `df.head()`

Out[7]:

	stay	age	hmo	died
0	4	4	0	0.0

```
df.info()
```

1	9	4	1	0.0
2	3	7	1	1.0
3	9	6	0	0.0
4	1	7	0	1.0

In [8]:

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1493 entries, 0 to 1492 Data  
columns (total 4 columns):  
#   Column  Non-Null Count  Dtype  
---  ---  
0  stay    1493 non-null   int64  
1  age     1493 non-null   int64  
2  hmo     1493 non-null   int64  
3  died    1493 non-null   float64 dtypes: float64(1), int64(3) memory usage: 46.8 KB
```

```
In [9]: In #Convert Variables
df['hmo'] = df['hmo'].astype('category') df['died'] =
df['died'].astype('category')
```

```
[10]: #Check Overdispersion print(
      "Mean:",
      df['stay'].mean()
    )

    print(
      "Variance:", df['stay'].var()
    )

    #Variance >> Mean so overdispersion exists
```

Mean: 9.728734092431345

```
#Check Zero Counts print(
  "Number of zeros:",
  (df['stay'] == 0).sum() )
```

Variance: 66.14419390578749

```
In [11]:
```

Number of zeros: 0

In [12]:

```
#Prepare Variables #Response

variable:

y = df['stay']

#Predictors:

X = pd.get_dummies( df[['age', 'hmo',
                        'died']],
                    drop_first=True
) #Convert to numeric:

X = X.astype(float)

#Add constant:

X = sm.add_constant(X)
```

In [13]:

```
#Fit Zero-Truncated Negative Binomial Model

#Model:stay~age+hmo+died
ztnb_model = TruncatedLFNegativeBinomialP(
    endog=y, exog=X
).fit()
```

Optimization terminated successfully.

Current function value: 3.185050

Iterations: 16

Function evaluations: 18

Gradient evaluations: 18

In [14]:

```
print(ztnb_model.summary())
```

TruncatedLFNegativeBinomialP Regression Results

=====

```
===== Dep. Variable:          stay No. Observations:
1493 Model:          TruncatedLFNegativeBinomialP Df Residuals:
1488 Method:          MLE Df Model:
3 Date:              Sun, 10 May 2026 Pseudo R-squ.:
0.003263 Time:        13:06:20 Log-Likelihood:
-4755.3 converged:    True LL-Null: -4770.8 Covariance
Type:                nonrobust LLR p-value:
7.955e-07
```

=====

	coef	std err	z	P> z	[0.025	0.975]
--	------	---------	---	------	--------	--------

-----const							2.4083	0.072	
33.457	0.000	2.267	2.549	age	-0.0157	0.013	-1.197	0.231	-0.041
0.010	hmo_1	-0.1471	0.059	-2.483	0.013	-0.263	-0.031	died_1.0	-0.2178
0.046	-4.718	0.000	-0.308	-0.127	alpha	0.5663	0.031	18.132	0.000
0.505									0.628
=====									
=====									

- $\text{Alpha} = 0.5663$
- $\text{P-value} < 0.001$

The significant positive alpha value indicates overdispersion exists in the data. This supports the use of the Negative Binomial model instead of the Poisson model.

Overall Conclusion: The analysis suggests that insurance type and mortality status significantly affect hospital length of stay, while age does not show a statistically significant effect. The Zero-Truncated Negative Binomial model is appropriate because the data are overdispersed and exclude zero-day hospital stays. """)

In []: